

# Allocating Memory

# Kernel Allocate Memory API

➤ **kmalloc()**

➤ **kfree**

➤ **\_\_get\_free\_page()**

➤ **\_\_free\_page()**

➤ **\_\_get\_free\_pages()**

➤ **\_\_free\_pages()**

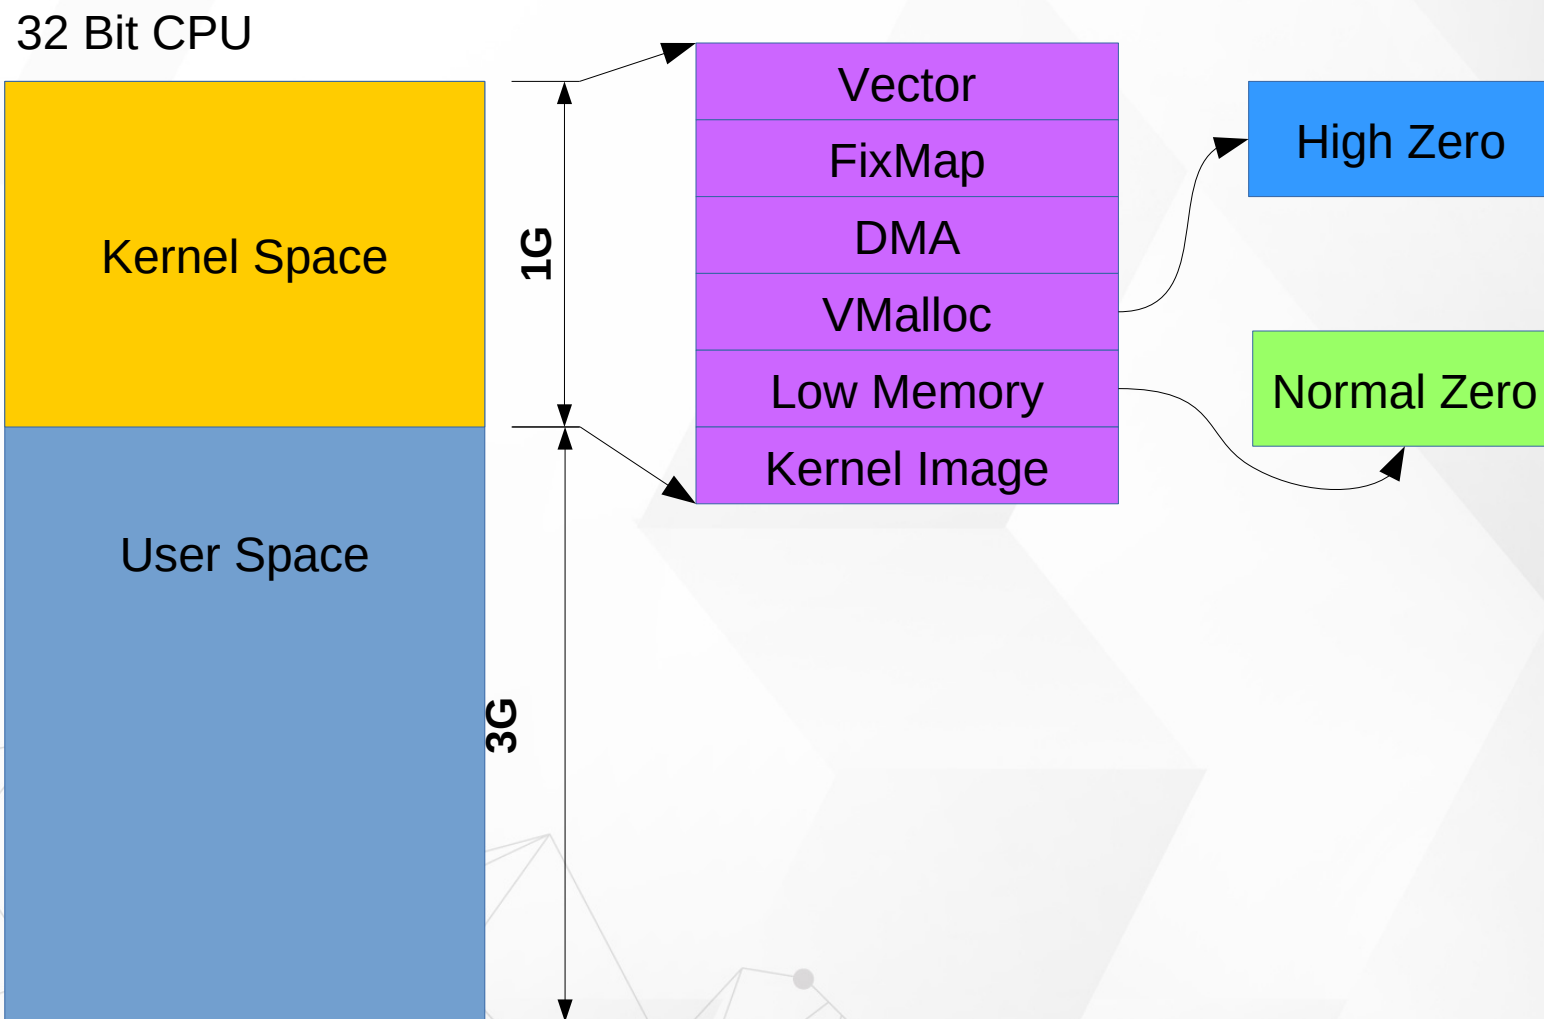
➤ **vmalloc()**

➤ **vfree()**

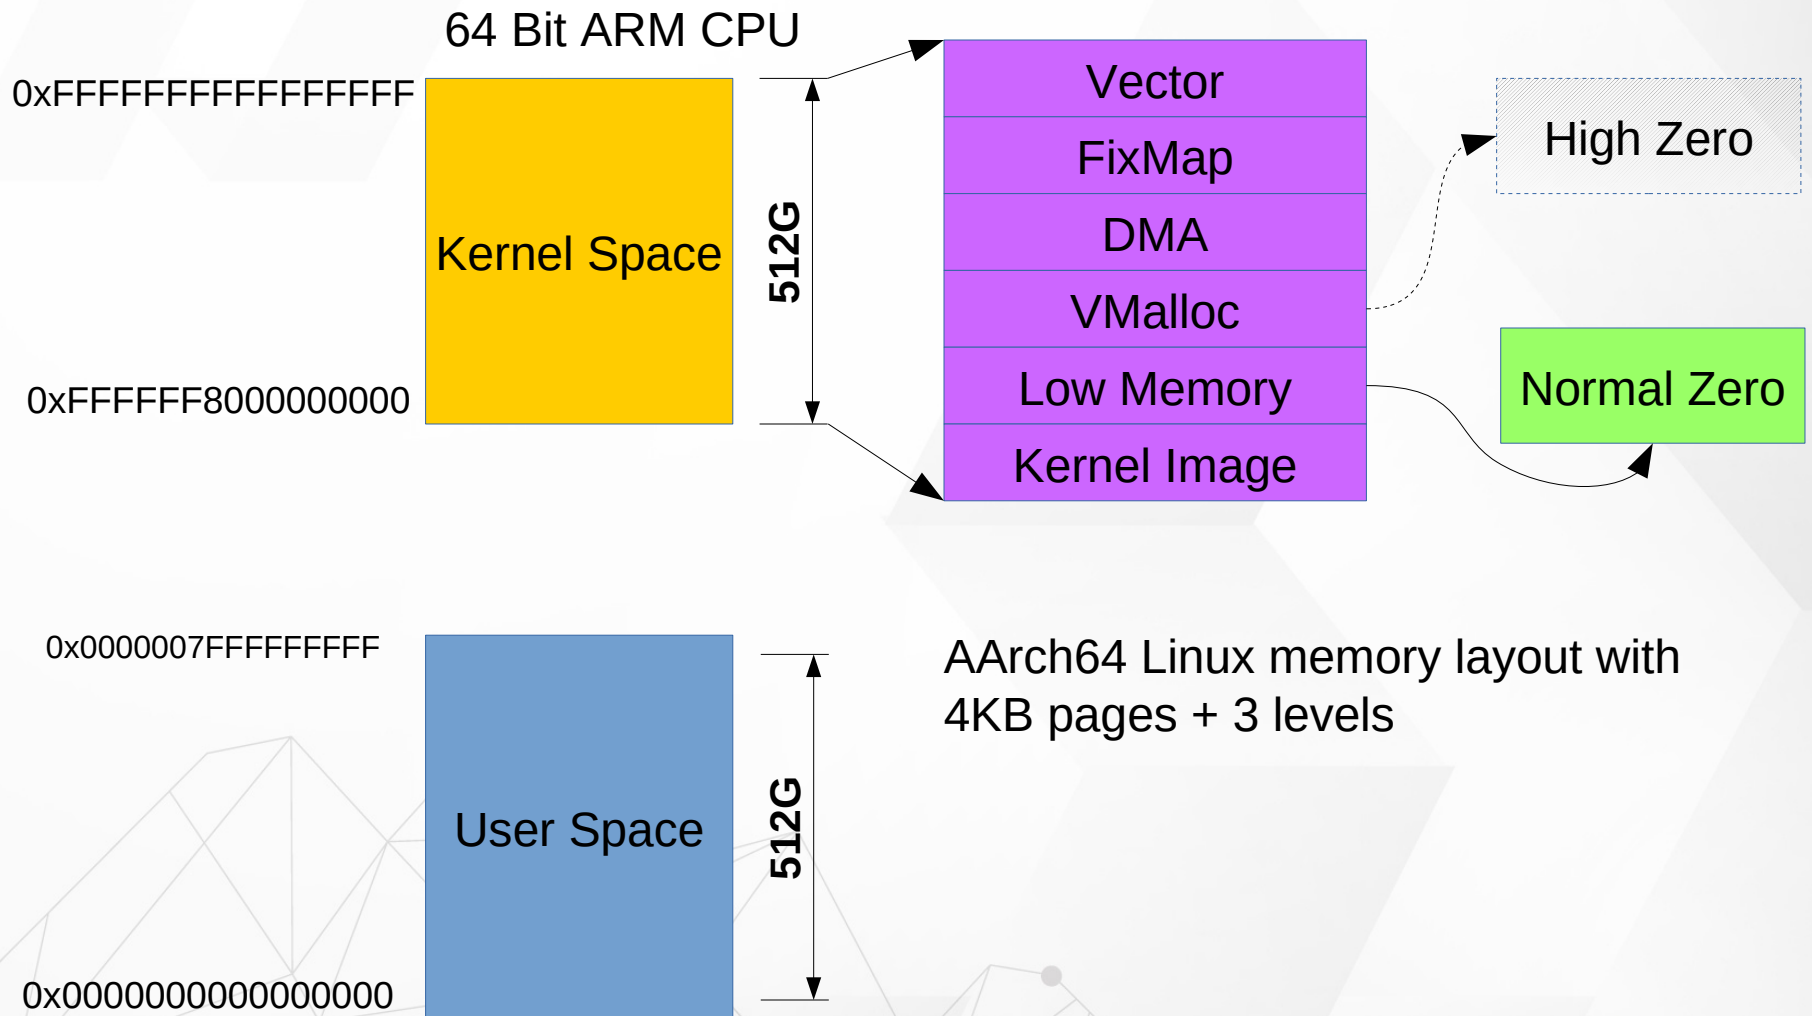
➤ **ioremap()**

➤ **iounmap()**

# Kernel Memory Configuration



# Kernel Memory Configuration



# The Flags Argument

## » **kmalloc** prototype

```
#include <linux/slab_def.h>
void *kmalloc(size_t size, int flags);
```

## » **Flags**

### » **GFP\_KERNEL**

- the most commonly used flag
- **Cannot** be used in **atomic** context

# kmalloc

➤ `kmalloc(size_t size, gfp_t flags);`

➤ Allocates **consecutive** virtual/physical memory pages

➤ Offset by **PAGE\_OFFSET**

- **PAGE\_OFFSET 2<sup>12</sup> (4K)**

➤ No changes to page tables

➤ Allocate memory limit : **128K**

# Memory Zones

- Normal memory

  - **GFP\_KERNEL**

- DMA-capable memory

  - **\_\_GFP\_DMA**

  - Platform dependent

  - First 16MB of RAM on the x86 for ISA devices

  - PCI devices have no such limit

- High memory

  - **\_\_GFP\_HIGHMEM**

  - Platform dependent

  - > 32-bit addressable range

# Driver SHT21

```
static int sht21_probe(struct i2c_client *client,
                      const struct i2c_device_id *id)
{
    struct device *dev = &client->dev;
    struct device *hwmon_dev;
    struct sht21 *sht21;

    if (!i2c_check_functionality(client->adapter,
                                I2C_FUNC_SMBUS_WORD_DATA)) {
        dev_err(&client->dev,
                "adapter does not support SMBus word transactions\n");
        return -ENODEV;
    }

    sht21 = devm_kzalloc(dev, sizeof(*sht21), GFP_KERNEL);
    if (!sht21)
        return -ENOMEM;

    sht21->client = client;

    mutex_init(&sht21->lock);

    hwmon_dev = devm_hwmon_device_register_with_groups(dev, client->name,
                                                         sht21, sht21_groups);
    return PTR_ERR_OR_ZERO(hwmon_dev);
}
```

kmalloc and flag



# The Size Argument

- › Kernel manages physical memory in **pages**
  - › Needs special management to allocate small memory chunks
- › Linux creates pools of memory objects in fixed sizes
  - › 32-byte
  - › 64-byte,
  - › 128-byte memory objects
- › **kmalloc create smallest allocation unit**
  - › 32 or 64 bytes
- › Largest portable allocation unit is **128KB**

# Supplement

# get\_free\_page and Friends

➤ For allocating **big chunks of memory**, it is more efficient to use a **page-oriented** allocator

➤ To allocate pages, call

Kernel API

```
/* returns a pointer to a zeroed page */
get_zeroed_page(unsigned int flags);

/* does not clear the page */
__get_free_page(unsigned int flags);

/* allocates multiple physically contiguous pages */
__get_free_pages(unsigned int flags, unsigned int order);
```

# get\_free\_page and Friends

## » flags

» Same as **flags** for **kmalloc**

## » order

» Allocate  $2^{\text{order}}$  pages

- **order** = 0 for 1 page
- **order** = 3 for 8 pages

» Can use **get\_order(size)** to find out **order**

» Maximum allowed value is about 10 or 11

» See **/proc/buddyinfo** statistics

# Page and order

Order

$2^{10}$

1024 Pages

$2^9$

512 Pages

$2^8$

256 Pages

$2^7$

128 Pages

$2^6$

64 Pages

$2^5$

32 Pages

$2^4$

16 Pages

$2^3$

8 Pages

$2^2$

4 Pages

$2^1$

2 Pages

$2^0$

1 Page

# get\_free\_page and Friends

- Subject to the same rules as **kmalloc**
- To free pages, call

```
void free_page(unsigned long addr);  
void free_pages(unsigned long addr, unsigned long order);
```

- Make sure to free the same number of pages
  - Or the memory map becomes corrupted

# Sample – get\_free\_page (1)

## 3-1-2-buddy

```
static int buddy_init(void)
{
    printk("Buddy driver installed...\n");

    order = 4;

    ptr = (void *)__get_free_pages(GFP_KERNEL, order);
    if (ptr)
        printk("2^%d * %ld bytes allocated\n", order, PAGE_SIZE);

    order2 = get_order(PAGE_SIZE * 4);
    ptr2 = (void *)__get_free_pages(GFP_KERNEL, order2);
    if (ptr2)
        printk("2^%d * %ld bytes allocated\n", order2, PAGE_SIZE);

    return 0;
}
```

# Sample – get\_free\_page (2)

```
static void buddy_exit(void)
{
    printk("buddy driver removed...\n");
    printk("buddy driver removed order  %d ...\n", order);
    printk("buddy driver removed order2  %d ...\n", order2);

    free_pages((unsigned long)ptr, order);
    free_pages((unsigned long)ptr2, order2);
}
```



# vmalloc and Friends

- » Allocates a **virtually contiguous** memory region
  - » Not consecutive pages in physical memory
    - Each page retrieved with a separate **alloc\_page** call
    - Less efficient
  - » Can sleep (cannot be used in atomic context)

# vmalloc and Friends

## » vmalloc-related prototypes

```
#include <linux/vmalloc.h>

void *vmalloc(unsigned long size);
void vfree(void * addr);
void *ioremap(unsigned long offset, unsigned long size);
void iounmap(void * addr);
```

# vmalloc and Friends

- » **ioremap** builds page tables
  - » **Does not allocate memory**
  - » Takes a physical address (**offset**) and return a virtual address
    - Useful to map the address of a Hardware buffer to kernel space
  - » Should use **readb** and other relation functions to access remapped memory